How Much Data is Enough? Benchmarking Transfer Learning for Few Shot ECG Image Classification

Sathvik Bhaskarpandit[1] [a]

[1]Department of Computer Science and Information Systems, BITS Pilani Hyderabad Campus, Hyderabad, India
f20191200@hyderabad.bits-pilani.ac.in

Abstract: Over the past couple of decades, numerous research works have been conducted to study and detect abnormalities from ECG signals. In this direction, several deep learning models have been proposed to detect these abnormalities and aid healthcare experts in their diagnoses. Although many of these deep learning approaches utilize ECG signals as input, only a handful use images of patients' ECGs themselves, that are often stored in hospitals and diagnostic centres. This work aims to study ECG images under the few-shot learning scenario. More specifically, it aims to study the effectiveness of transfer learning for few-shot ECG image classification, and how classification performance varies with the amount of training data available. Results show that models such as ResNet and EfficientNet are able to classify images with great success with around 20 images per class, with accuracy even crossing 99.5%. Yet under extreme data unavailability cases, such as 5-shot learning and lower, transfer learning proves to be unreliable to be put to use in healthcare for automated classification of ECG images.

1 INTRODUCTION

An electrocardiogram (ECG) is an electrical recording of the heart representing the cardiac cycle, on a graph representing the electrical activity of the heart obtained by connecting electrodes adapted to the body surface. It is a widely used noninvasive medical test used for measuring the heart condition by tracking the heart's electrical activity. It plays a huge role in the field of medicine and healthcare, ranging from detection of cardiac diseases to vascular diseases to COVID-19. ECG contains plenty of information that directly reflects cardiac physiology since its morphological and temporal features are produced from cardiac electrical and structural variations. The waves produced by ECG signals are characterized by their shapes and durations. When certain changes affect certain characteristics of these waves, a heart defect is considered to be present.

While an experienced cardiologist can distinguish different types of cardiology abnormalities by visually referencing the ECG waveform pattern, a machine learning (ML) approach can improve the diagnostic efficiency. Therefore, detection and treatment of anomalies have become the main research topics in the field of cardiac care and the information processing domain. Numerous methods have been proposed to classify, as well as automatically detect various types of abnormalities from ECG signals. Early methods include use of recursive filters (Zeraatkar et al., 2011) and wavelet transforms (Addison, 2005) to detect arrhythmia. Then, various feature extraction techniques such as peak detection (Khazaee, 2013), QRS complex detection (Li et al., 2016), RR interval analysis (Tsipouras et al., 2005), Empirical Mode Decomposition (Izci et al., 2018), etc. were employed to aid in classification. Machine learning models such as support vector machine (Asl et al., 2008), logistic regression (Behadada et al., 2016), etc. were used to classify these signals from the hand-crafted features. However, with the advent of deep learning (DL), DL models began to predominate as they could automatically extract complex features. Popular DL models include the use of multi layer perceptrons (MLP) (Savalia and Emamian, 2018), convolutional neural networks (CNN) (Wu et al., 2021), long-short term memory (LSTM) (Gao et al., 2019) and deep belief networks (DBN) (Gourisaria et al., 2021) .

While ECG signals themselves are an invaluable source of data, a considerable amount of ECG data is stored in hospitals in the form of images. Extracting patterns and classifying these images proves to be a

[a] https://orcid.org/0000-0001-6201-4975

difficult task. Several studies have been carried out on classification from signals, whereas far, far fewer studies have been carried out for classification from images. These methods that classify ECG images utilize DL approaches to extract complex features from the images and subsequently perform classification (Mohamed et al., 2015; Jun et al., 2018).

However, DL models suffer from a fundamental drawback: They require a large number of training examples to achieve satisfactory performance. While humans are able to identify objects by simply looking at a couple of such instances, DL models are unable to do the same. Further, such a large number of images in places such as hospitals may not be labelled initially. A doctor or subject matter expert (SME) is required to label the data (Gupta et al., 2021). Manual labelling may lead to inexact and noisy labels. The cost and time required for labelling such a large number of training examples is high and not scalable. Therefore, there is a need for models that can perform classification with high accuracy, but with a limited number of labelled training images. This is where few-shot learning (FSL) comes into the picture.

FSL aims to classify a set of testing examples, known as the query set, given a limited number of training examples, known as the support set. Transfer learning (TL) is a popular approach for FSL. It involves pre-training a model on a large dataset, fine-tuning the model on the support set, then finally testing its performance on the query set (Weimann and Conrad, 2021; Venton et al., 2020; Salem et al., 2018).

This work aims to conduct experiments to study the effectiveness of transfer learning in few-shot ECG image classification (FSEIC). In particular, several popular pretrained image classification architectures are used and fine-tuned on the ECG images. The effect of amount of data available during training on the classification performance is observed. An estimate of the minimum number of labelled images required to achieve a high accuracy is calculated.

2 EXPERIMENTAL SETUP

2.1 Dataset and Data Preprocessing

The data used for fine-tuning is taken from the ECG Images dataset of Cardiac and COVID-19 Patients (Khan et al., 2021). It consists of ECG images collected from different health care institutes across Pakistan. The ECG signals in the images themselves are sampled at 500 Hz. All the collected data was manually reviewed and labelled by Senior Medical Professionals. Each of the images belongs to one of three classes: normal patient, patient with abnormal heartbeat, or patient with myocardial infarction, with each class having approximately 250 images.

Each image contains metadata such as patient name, ID, height, weight, time of recording, etc. Such details are cropped out of the images to retain only a grid with the snapshots of the ECG signal recordings. The images are then resized to a size of (128, 128, 3) using image anti-aliasing. Finally, depending on the pretrained model used, additional preprocessing is done to make the images ready to be inputted into the model.

2.2 Sampling of Tasks

After the images have been preprocessed, FSL tasks are sampled from the images. Each task consists of randomly sampling k images per class as the support set, fine-tuning the model on this support set, then testing the model on another set of randomly sampled images, i.e. q images per class, known as the query set. The support set and query set are mutually disjoint.

2.3 Model Fine-Tuning

Each base model is initialized with the weights of a popular image-classification model pretrained on the ImageNet dataset (Deng et al., 2009). In particular, the base models used are convolutional neural networks that include VGG16 (Simonyan and Zisserman, 2014), DenseNet121 (Huang et al., 2017), InceptionV3 (Szegedy et al., 2015), ResNet101 (He et al., 2015), and EfficientNet models B0 to B7 (Tan and Le, 2019). Any fully connected layers connected to the top of the base model are removed. The weights of the layers of the base model are frozen. Therefore, the base model acts as a feature extractor for the images.

On top of this base model, a global average pooling (GAP) layer is added, followed by batch normalization and dropout. The output of the GAP layer is subsequently fed into a softmax layer to output the class probabilities. The trainable parameters therefore include those of the batch normalization layer and weights of the final softmax layer. Categorical crossentropy is taken as the loss function to be minimized while fine-tuning the model on the support set. The Adam optimizer (Kingma and Ba, 2014) with an initial learning rate of 10^{-2} and batch size of 16 is used to arrive at the optimal weights. Early stopping is used as the criterion for stopping training. The fine-tuned model is finally used to obtain class predictions for the query set. Metrics such as accuracy and F1

Table 1: Accuracy and F1 Scores of the various transfer learning models with different pretrained base models. Values are in percentage. Underlined values represent the model with the best accuracy or F1 score for the given value of k. The table headings B0 to B7 represent EfficientNet model variants.

k	Metric	VGG16	DenseNet	InceptionV3	ResNet	B0	B1	B2	B3	B4	B5	B6	B7
1	Accuracy	39.67	41.33	47.33	50.07	64.80	58.60	59.10	52.13	40.07	51.87	41.73	45.73
	F1	29.01	35.53	32.95	46.92	61.01	53.39	53.74	42.92	30.48	47.07	32.78	41.69
2	Accuracy	65.03	61.33	47.09	60.33	79.93	79.47	80.63	57.33	57.87	65.21	44.53	41.40
	F1	63.96	61.24	36.52	58.89	79.40	79.11	79.97	49.85	51.59	64.29	34.38	40.24
3	Accuracy	55.67	61.03	53.67	72.93	92.82	87.43	86.21	73.07	73.47	68.80	51.33	60.73
	F1	51.67	56.08	51.41	72.71	92.79	87.4	86.04	70.66	72.17	67.82	45.21	62.88
4	Accuracy	53.01	77.33	56.50	83.4	93.98	92.80	93.10	83.07	82.85	78.67	75.73	78.25
	F1	46.27	75.86	51.16	83.2	93.97	92.78	93.05	82.83	82.21	78.45	75.55	80.32
5	Accuracy	71.33	86.33	64.33	86.00	96.07	92.67	94.73	84.67	84.80	83.00	78.40	82.40
	F1	70.65	86.33	57.09	85.94	96.06	92.65	94.71	84.53	84.88	82.99	78.14	84.30
10	Accuracy	67.33	87.83	81.00	90.73	97.15	97.00	95.60	92.07	90.80	91.53	88.80	90.13
	F1	63.82	87.75	80.88	90.74	97.16	96.98	95.58	92.03	90.78	91.59	88.81	90.36
20	Accuracy	85.33	92.67	88.33	95.53	98.40	98.63	99.20	96.93	96.00	96.27	94.87	95.41
	F1	85.38	92.68	88.11	95.54	98.39	98.63	99.22	96.91	95.98	96.24	94.85	95.45
30	Accuracy	86.33	94.33	89.67	97.67	99.03	99.01	99.37	98.93	96.87	97.87	95.62	96.86
	F1	86.29	94.33	89.65	97.65	99.01	99.02	99.35	98.93	96.85	97.86	95.59	96.90
40	Accuracy	86.67	96.07	93.17	98.93	99.08	99.43	99.57	98.80	98.33	98.13	96.73	97.66
	F1	86.55	95.99	93.16	98.91	99.10	99.37	99.57	98.74	98.33	98.19	96.72	97.69
50	Accuracy	88.67	97.33	93.33	99.61	99.28	99.57	99.67	99.61	98.13	98.73	96.87	98.53
	F1	88.53	97.32	93.26	99.60	99.28	99.53	99.66	99.60	98.12	98.71	96.86	98.59

score are recorded for each task. Since we use an equal number of samples per class for the query set, the results are primarily discussed in terms of accuracy in section 3.

For each TL model, k is varied from 1 to 50 to observe the effect of amount of training data on the classification performance. The value of q is fixed at 100. For each value of k, 20 tasks are sampled, each time with a different random seed, and the final metrics reported are taken as the average of these 20 tasks. For fairness of comparison, the sampled tasks for a given value of k are the same across all the TL models.

3 RESULTS AND DISCUSSION

In this section, the results and plausible explanations for the recorded observations are presented. Table 1 shows the results of the various TL models for different values of k.

From the table we notice a couple of interesting observations. For $k = 1$, the accuracies and F1 scores of the models are quite low, with all being less than 65%. There is a relatively large increase in accuracy from $k = 1$ to $k = 2$, with a smaller increase from $k = 2$ to $k = 3$. From $k = 1$ till $k = 5$, with the exception of $k = 2$, EfficientNetB0 is the best performing model. For larger values of k, EfficientNetB2 seems to be the best performing model. If we take into the consideration the largest value of $k = 50$, only ResNet and EfficientNet models B0 to B3 cross 99%. VGG16 and

InceptionV3 models perform poorly, while DenseNet and EfficientNet models B4 to B7 show promising performance, at around 97-99%.

To have a deeper look at the effect of the amount of data during training, the accuracies of each model are plotted versus k. Figures 1 and 2 show the plots for $k = 1$ till $k = 19$, and $k = 20$ till $k = 50$, respectively. We observe that there is a sharper increase in accuracy in the beginning for smaller values of k, which becomes more gradual for later values of k. The increase in accuracy is not monotonic but irregular, primarily due to the scarcity of data and minorly due to the stochastic nature of optimization of neural networks. For almost all values of k, EfficientNet models B0-B2 show the best performance. We conclude that these models have the best ability to extract useful features from the images and classify them into one of the class labels. Until $k = 8$, the B0 model seems to perform better while the B2 model seems to be the best performing model for most values of k later. This is possibly due to the fact that more complex and deeper models initially overfit for smaller values of k as there is lesser amount of training data and tend to result in high-dimensional features from the GAP layer. As the amount of training data grows, more complex models learn better features than simpler ones and these high-dimensional features are useful in differentiating between the classes. However, EfficientNet models B5-B7 are still too complex for the amount of data being experimented with. In this regard, it is observed that ResNet performs poorly for smaller values of k but extremely well for larger values.

Figure 1: Plot of accuracy as a function of *k* for *k*=1 to *k*=20

Figure 2: Plot of accuracy as a function of *k* for *k*=21 to *k*=50

We also seek to answer the question: How much data is required to obtain a certain threshold accuracy δ that we consider is useful in a real world scenario, such as medical diagnoses in hospitals, and is it a reasonable value? Medical diagnoses carry a high degree of responsibility and medical research strives to make such diagnoses impeccable. Thus a high value of δ is obviously preferred. Several research works that have been conducted to classify ECG signals achieve classification accuracies of around 99.5% (Ji et al., 2019; Shoughi and Dowlatshahi, 2021) on the MIT-BIH database (Moody and Mark, 1992). We observe that for $\delta = 99\%$, $k_{min} = 24$ using EfficientNetB2, Similarly, for $\delta = 99.5\%$, $k_{min} = 40$ using Efficient-NetB2. Obtaining approximately 40 images per class for training a model is certainly possible, but would

take a reasonable amount of time. Additionally, it is relatively easier to obtain ECG images of normal patients than patients with myocardial infarction. In case the disease being observed is rarer, it may be even more difficult to collect such images. Considering that almost all models perform poorly given only a couple of training samples per class, it may not be a wise idea to rely on deep learning to automatically classify a patient's condition in such a scenario.

4 CONCLUSION AND FUTURE WORK

This work has shown that transfer learning using popular image classification architectures is a promising direction for few-shot ECG image classification. With around 20 images per class available for training, models such as ResNet and EfficientNet are able to achieve accuracies of at least 99%.

However, when the training set comprises 5 images per class or fewer, simple transfer learning fails to classify ECG images with high accuracy. In this direction, other algorithms that work well in low-shot and class imbalance scenarios, can be explored. Other few-shot learning methods can be compared with transfer learning and observed to see how well they perform with the amount of labelled data available. It is also worth experimenting with data from different ethnicities and regions as the current work deals with data taken from one region only.

REFERENCES

Addison, P. S. (2005). Wavelet transforms and the ecg: a review. *Physiological measurement*, 26(5):R155.

Asl, B. M., Setarehdan, S. K., and Mohebbi, M. (2008). Support vector machine-based arrhythmia classification using reduced features of heart rate variability signal. *Artificial intelligence in medicine*, 44(1):51–64.

Behadada, O., Trovati, M., Chikh, M. A., Bessis, N., and Korkontzelos, Y. (2016). Logistic regression multinomial for arrhythmia detection. In *2016 IEEE 1st International Workshops on Foundations and Applications of Self* Systems (FAS* W)*, pages 133–137. IEEE.

Deng, J., Dong, W., Socher, R., Li, L.-J., Li, K., and Fei-Fei, L. (2009). Imagenet: A large-scale hierarchical image database. In *2009 IEEE conference on computer vision and pattern recognition*, pages 248–255. Ieee.

Gao, J., Zhang, H., Lu, P., and Wang, Z. (2019). An effective lstm recurrent network to detect arrhythmia on imbalanced ecg dataset. *Journal of healthcare engineering*, 2019.

Gourisaria, M. K., Harshvardhan, G., Agrawal, R., Patra, S. S., Rautaray, S. S., and Pandey, M. (2021). Arrhythmia detection using deep belief network extracted features from ecg signals. *International Journal of E-Health and Medical Communications (IJEHMC)*, 12(6):1–24.

Gupta, P., Bhaskarpandit, S., and Gupta, M. (2021). Similarity learning based few shot learning for ecg time series classification. In *2021 Digital Image Computing: Techniques and Applications (DICTA)*, pages 1–8. IEEE.

He, K., Zhang, X., Ren, S., and Sun, J. (2015). Deep residual learning for image recognition. arxiv 2015. *arXiv preprint arXiv:1512.03385*.

Huang, G., Liu, Z., Van Der Maaten, L., and Weinberger, K. Q. (2017). Densely connected convolutional networks. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 4700–4708.

Izci, E., Ozdemir, M. A., Sadighzadeh, R., and Akan, A. (2018). Arrhythmia detection on ecg signals by using empirical mode decomposition. In *2018 Medical Technologies National Congress (TIPTEKNO)*, pages 1–4. IEEE.

Ji, Y., Zhang, S., and Xiao, W. (2019). Electrocardiogram classification based on faster regions with convolutional neural network. *Sensors*, 19(11):2558.

Jun, T. J., Nguyen, H. M., Kang, D., Kim, D., Kim, D., and Kim, Y.-H. (2018). Ecg arrhythmia classification using a 2-d convolutional neural network. *arXiv preprint arXiv:1804.06812*.

Khan, A. H., Hussain, M., and Malik, M. K. (2021). Ecg images dataset of cardiac and covid-19 patients. *Data in Brief*, 34:106762.

Khazaee, A. (2013). Heart beat classification using particle swarm optimization. *International Journal of Intelligent Systems and Applications*, 5(6):25.

Kingma, D. P. and Ba, J. (2014). Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.

Li, P., Wang, Y., He, J., Wang, L., Tian, Y., Zhou, T.-s., Li, T., and Li, J.-s. (2016). High-performance personalized heartbeat classification model for long-term ecg signal. *IEEE Transactions on Biomedical Engineering*, 64(1):78–86.

Mohamed, B., Issam, A., Mohamed, A., and Abdellatif, B. (2015). Ecg image classification in real time based on the haar-like features and artificial neural networks. *Procedia Computer Science*, 73:32–39.

Moody, G. B. and Mark, R. G. (1992). Mit-bih arrhythmia database.

Salem, M., Taheri, S., and Yuan, J.-S. (2018). Ecg arrhythmia classification using transfer learning from 2-dimensional deep cnn features. In *2018 IEEE biomedical circuits and systems conference (BioCAS)*, pages 1–4. IEEE.

Savalia, S. and Emamian, V. (2018). Cardiac arrhythmia classification by multi-layer perceptron and convolution neural networks. *Bioengineering*, 5(2):35.

Shoughi, A. and Dowlatshahi, M. B. (2021). A practical system based on cnn-blstm network for accurate classification of ecg heartbeats of mit-bih imbalanced dataset. In *2021 26th International Computer Conference, Computer Society of Iran (CSICC)*, pages 1–6. IEEE.

Simonyan, K. and Zisserman, A. (2014). Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*.

Szegedy, C., Vanhoucke, V., Ioffe, S., Shlens, J., and Wojna, Z. (2015). Rethinking the inception archi-

tecture for computer vision. 2015. *arXiv preprint arXiv:1512.00567*.

Tan, M. and Le, Q. (2019). Efficientnet: Rethinking model scaling for convolutional neural networks. In *International conference on machine learning*, pages 6105–6114. PMLR.

Tsipouras, M. G., Fotiadis, D. I., and Sideris, D. (2005). An arrhythmia classification system based on the rr-interval signal. *Artificial intelligence in medicine*, 33(3):237–250.

Venton, J., Aston, P. J., Smith, N. A., and Harris, P. M. (2020). Signal to image to classification: Transfer learning for ecg. In *2020 11th Conference of the European Study Group on Cardiovascular Oscillations (ESGCO)*, pages 1–2. IEEE.

Weimann, K. and Conrad, T. O. (2021). Transfer learning for ecg classification. *Scientific reports*, 11(1):1–12.

Wu, M., Lu, Y., Yang, W., and Wong, S. Y. (2021). A study on arrhythmia via ecg signal classification using the convolutional neural network. *Frontiers in computational neuroscience*, page 106.

Zeraatkar, E., Kermani, S., Mehridehnavi, A., Aminzadeh, A., Zeraatkar, E., and Sanei, H. (2011). Arrhythmia detection based on morphological and time-frequency features of t-wave in electrocardiogram. *Journal of Medical Signals and Sensors*, 1(2):99.